

```

/
*****
*
* Name: Roger Shore
* Date: 9-8-2016
* Class: CSC4310
* Location: ~/HPU/csc4310/mpi
*
* FILE: sum_array.c
* DESCRIPTION:
*   MPI Example - Array Assignment - C Version
*   This program demonstrates a simple data decomposition. The master task
*   first initializes an array and then distributes an equal portion that
*   array to the other tasks. After the other tasks receive their portion
*   of the array, they perform an addition operation to each array element.
*   They also maintain a sum for their portion of the array. The master task
*   does likewise with its portion of the array. As each of the non-master
*   tasks finish, they send their sum to the master which computes a total.
*
*   An MPI collective communication call is used to collect the sums
*   maintained by each task. Finally, the master task displays selected
*   parts of the final array and the global sum of all array elements.
*   NOTE: the number of MPI tasks must be evenly divided by 4.
*
* NOTE: This is an example. If it was code to be released, only
* the master node would output any values. The output from the other
* nodes is to verify the data flow and my understanding of the events
* that take place when the program is executed.
*
* To Compile:
*
* mpicc sum_array.c -o sum_array
*
* To execute:
*
* mpiexec --mca btl_tcp_if_include <ethernet_ID> -n 4 -hostfile actHosts
sum_array
*   either ifconfig or ip a to determine the ethernet ID for the machine.
* mpiexec --mca btl_tcp_if_include enp3s -n 4 -hostfile actHosts sum_array
*
*   NOTE: the number of MPI tasks must be evenly divided by 4.
*****/

#include "mpi.h"
#include <stdio.h>
#include <stdlib.h>
#include <sys/time.h>
#define ARRAYSIZE 1600
#define MASTER 0
#define TAG1 0
#define TAG2 1
#define TAG3 2

void cluster_node_process(int chunksize, int taskid);

```

```

void master_node_process(int chunksize, int numtasks, int taskid);
double findSum(double data[], int myoffset, int chunk, int myid);
double initArray(double data[], int n);
void sampleOutput(double data[], int chunksize, int numtasks);

int main (int argc, char *argv[])
{
    int numtasks, taskid, rc=1;
    int chunksize;

    /*** Initializations ***/
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numtasks);
    if (numtasks % 4 != 0) {
        printf("Quitting. Number of MPI tasks must be divisible by 4.\n");
        MPI_Abort(MPI_COMM_WORLD, rc);
        exit(0);
    }

    MPI_Comm_rank(MPI_COMM_WORLD, &taskid);

    printf ("MPI task %d has started...\n", taskid);
    chunksize = (ARRAYSIZE / numtasks);

    /*** Master task only ***/
    if (taskid == MASTER){
        master_node_process(chunksize, numtasks, taskid);
    }

    /*** Non-master tasks only ***/
    if (taskid > MASTER) {
        cluster_node_process(chunksize, taskid);
    } /* end of non-master */

    MPI_Finalize();
} /* end of main */

/* master_node_process - init an array of data to process
 *
 * Send on portion of the array to
 *
 * other cluster nodes. Keep one chunk
 *
 * for the master
 *
 * Find the sum of the data
 *
 * Receive the results from each node.
 * precondition: chunksize, numtasks, and taskid are established
 * in the main for all processes
 * postcondition: nothing is returned from the function.
 */
void master_node_process(int chunksize, int numtasks, int taskid)
{

```

```

double data[ARRAYSIZE];
int offset, source;
struct timeval startTime, stopTime;
double start, stop, diff;
double mysum, sum;
double t1,t2,result;
MPI_Status status;
double sum2,tmpsum;
gettimeofday(&startTime,NULL);
t1 = MPI_Wtime();

sum = initArray(data,ARRAYSIZE);
printf("What we should get - Initialized array sum = %e\n",sum);

/* Send each task its portion of the array - master keeps 1st part */
offset = chunksize;
for (int dest=1; dest<numtasks; dest++) {
    MPI_Send(&offset, 1, MPI_INT, dest, TAG1, MPI_COMM_WORLD);
    MPI_Send(&data[offset], chunksize, MPI_DOUBLE, dest, TAG2,
MPI_COMM_WORLD);
    printf("Sent %d elements to task %d offset= %d\n",
chunksize,dest,offset);
    offset = offset + chunksize;
}

/* Master does its part of the work */
offset = 0;
mysum = findSum(data, offset, chunksize, taskid);
sum2 = mysum;

/* Wait to receive results from each task */
/* Get final sum and print sample results */
for (int i=1; i<numtasks; i++) {
    source = i;
    MPI_Recv(&offset, 1, MPI_INT, source, TAG1, MPI_COMM_WORLD, &status);
    MPI_Recv(&data[offset], chunksize, MPI_DOUBLE, source, TAG2,
MPI_COMM_WORLD, &status);
    MPI_Recv(&tmpsum, 1, MPI_DOUBLE, source, TAG3,
MPI_COMM_WORLD, &status);
    sum2 += tmpsum;
}

sampleOutput(data,chunksize,numtasks);
printf("*** What we got - Final sum= %e ***\n",sum2);

gettimeofday(&stopTime,NULL);
t2 = MPI_Wtime();
start = startTime.tv_sec + (startTime.tv_usec/1000000.0);
stop = stopTime.tv_sec + (stopTime.tv_usec/1000000.0);

diff = stop - start;
result = t2-t1;

printf("Time %f\n",diff);

```

```

    printf("MPI_Wtime %f\n",result);
}

/* cluster_node_process - receive a chunk of data from the master
 *                          find the sum of the data
 *                          send back the data and sum.
 * NOTE: there is no reason to send the data back, just demo
 *       how to send the array back.
 * precondition: chunksize and taskid are established in the main
 *              for all processes
 * postcondition: nothing is returned from the function.
 */
void cluster_node_process(int chunksize, int taskid)
{
    int source, dest, offset;
    MPI_Status status;
    double mysum;
    double data[ARRAYSIZE];

    /* Receive my portion of array from the master task */
    source = MASTER;
    MPI_Recv(&offset, 1, MPI_INT, source, TAG1, MPI_COMM_WORLD, &status);
    MPI_Recv(&data[offset], chunksize, MPI_DOUBLE, source, TAG2,
MPI_COMM_WORLD, &status);

    mysum = findSum(data, offset, chunksize, taskid);

    /* Send my results back to the master task */
    dest = MASTER;
    MPI_Send(&offset, 1, MPI_INT, dest, TAG1, MPI_COMM_WORLD);
    MPI_Send(&data[offset], chunksize, MPI_DOUBLE, MASTER, TAG2,
MPI_COMM_WORLD);
    MPI_Send(&mysum, 1, MPI_DOUBLE, MASTER, TAG3, MPI_COMM_WORLD);

    printf("\t\t\t\t\t*** Final sum= %e ***\n",mysum);
}

/* initArray - initialize an array with increasing values of i starting at 0
 * precondition: The data array is empty and is large enough to hold n ints
 * postcondition: the data array is updated.
 */
double initArray(double data[], int n)
{
    /* Initialize the array */
    double sum = 0;
    for(int i=0; i<n; i++) {
        data[i] = i * 1.0;
        sum = sum + data[i];
    }
    return sum;
}

/* findSum - find the sum of the elements within the chunk of data
 * precondition: The data array is loaded iio a specific chunk

```

```

    * precondition: the data array is updated.
    */
double findSum(double data[], int myoffset, int chunk, int myid)
{
    int i;
    double mysum;
    mysum = 0;
    for(i=myoffset; i < myoffset + chunk; i++) {
        mysum = mysum + data[i];
    }
    printf("\t\t\t\tTask %d mysum = %e\n",myid,mysum);
    return(mysum);
}

/* update - output the first 5 numbers each chunk of data
           passed to the non-master processes
    * precondition: The data array contains a collection of doubles
    *                 chunksize is the number of doubles in each set
    *                 of values. numtasks is the number of tasks
    *                 (or number of subsets) that process the data.
    * postcondition: no updates or changes in the data
    */
void sampleOutput(double data[],int chunksize, int numtasks)
{
    printf("Sample results: \n");
    int offset = 0;
    for (int i=0; i<numtasks; i++) {
        for (int j=0; j<5; j++)
            printf(" %e",data[offset+j]);
        printf("\n");
        offset = offset + chunksize;
    }
}

```